MILESTONE 3 - NARRATIVE FLOW CO WRITER
MODULE 3 – GUARD RAILS

1.Introduction -What is a guardrail?

The **Guardrails & Response Governance Engine** is a control module within the **Narrative Flow Co-Writer system** designed to strictly regulate how the AI generates responses. Its primary objective is to enforce well-defined boundaries so that the AI produces **only story-related creative content** and avoids responding to topics outside the narrative domain. Through these guardrails, the system ensures that the AI does not generate sexual content, political opinions or analysis, sensitive personal information, technical code, legal, medical, or financial advice, or any form of instructional or procedural guidance unrelated to storytelling. By implementing these restrictions, the module effectively transforms the AI from a **general-purpose conversational assistant into a specialized narrative-focused creative engine**, ensuring safe, controlled, and context-appropriate responses within the storytelling environment

---

2.Architecture

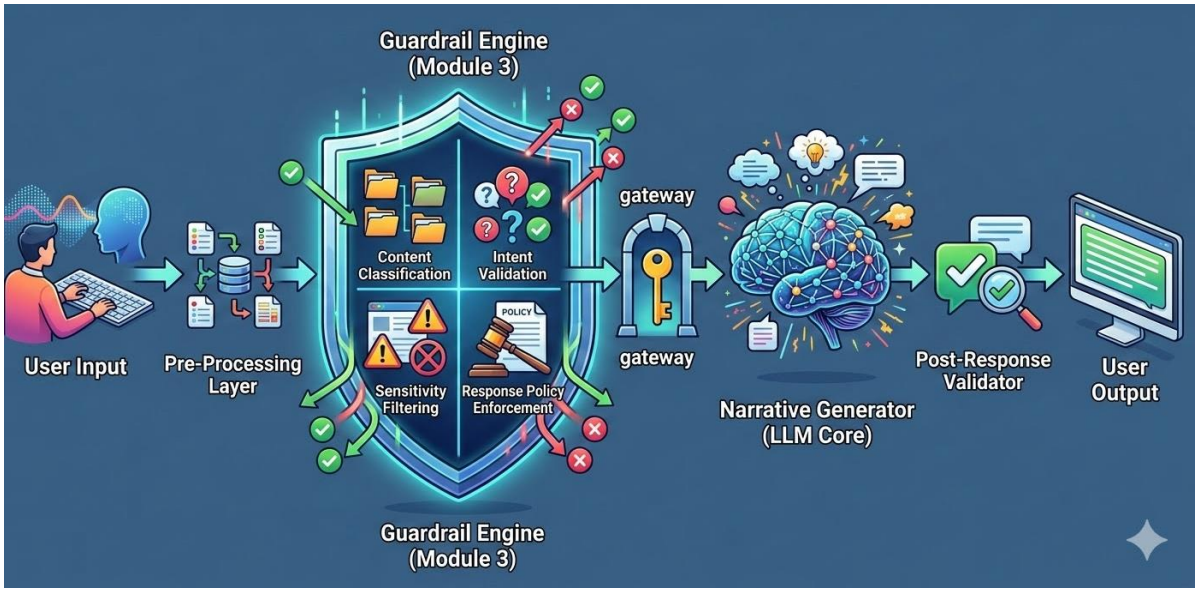

---

3. Functional Objectives

Function	Description
Intent Filtering	Detect whether the user request relates to story writing

Function	Description
Topic Restriction	Block restricted domains (sexual, political, etc.)
Output Constraining	Force narrative-only generation style
Refusal Formatting	Return structured denial when violation occurs
Continuous Monitoring	Validate both input and output

---

4. Restricted Content Domains

The AI must **hard-reject** the following categories:

4.1 Sexual Content

- Explicit content
- Adult material
- Suggestive narratives
- Minor-related content

4.2 Political Content

- Opinions on governments
- Policy arguments
- Political persuasion
- Ideological debates

4.3 Sensitive Information

- Personal data requests
- Identity details
- Confidential or private information

4.4 Code & Technical Assistance

- Programming code
- Debugging
- Technical system configuration
- Engineering tutorials

4.5 Professional Advice

- Legal advice
- Medical advice
- Financial advice
- Investment strategy
- Career consultation

---

5. Allowed Content Scope

The system is strictly limited to:

- Fiction writing
- Plot development
- Character arcs
- World-building
- Dialogue refinement
- Scene structuring
- Literary enhancement
- Tone adjustment
- Story pacing analysis

6.1. Enforcement Mechanism

Use a lightweight intent classifier:

- Story-related
- Technical request
- Political request
- Sexual content
- Advice-seeking
- Other restricted

If classification ≠ story → reject.

6.2 Rule-Based Filtering Layer

Example policy structure:

IF request contains:

- political entities
- sexual keywords
- code syntax markers
- advisory verbs ("should I invest", "what medicine")

THEN:

Reject request

6.3 Output Constraining Prompt

System instruction to LLM:

"You are a narrative co-writer. You must only generate fictional story content. If the request is not related to storytelling, you must refuse."

---

7. Refusal Response Standard

Refusals must:

- Be polite
- Be neutral
- Not explain internal policy logic
- Not suggest alternative restricted topics

Standard Refusal Template:

"I'm designed specifically to assist with creative storytelling. I'm unable to help with that request. If you'd like to continue working on a story, I'd be happy to help."

---

8. Edge Case Handling

Scenario	Action
Political content inside fictional world	Allowed if clearly fictional
Romantic subplot	Allowed if non-explicit
War in fantasy setting	Allowed
Real-world political leader in story	Block
Writing a fictional hacker	Allowed
Asking for actual hacking method	Block

9.Conclusion

- The Guardrails & Response Governance Engine establishes a compliance framework that ensures the AI operates with ethical constraints and remains strictly focused on storytelling tasks.
- It converts the system from a general-purpose chatbot into a specialized narrative engine, limiting responses to creative writing and story development.
- The module promotes ethical AI usage by preventing generation of restricted content such as political opinions, explicit sexual material, professional advice, or technical instructions.
- It enables creative domain specialization, ensuring the AI concentrates only on narrative construction, plot design, character development, and storytelling assistance.

- The system reduces liability and misuse risks by implementing strict response boundaries and rejecting requests outside the permitted domain.
- It improves brand clarity and product positioning by defining the AI clearly as a narrative-focused creative assistant.
- The module preserves content integrity by maintaining consistent storytelling-focused outputs without deviation into unrelated domains.
- It supports regulatory and compliance alignment by preventing generation of sensitive or advisory information that could create legal exposure.
- The guardrails enforce predictable system behaviour, where storytelling requests are supported while political, technical, or advisory queries are refused.
- For production deployment, the module uses layered enforcement mechanisms such as moderation filtering, keyword blacklists, semantic classifiers (e.g., BERT), and post-generation scanning to ensure all outputs remain compliant with narrative-only policies.

By
R.V. Aswath Narayanan
Aswathnarayanan48@gmail.com